

# Authentication and Authorization Patterns for Agents and MCP

Christian Posta - Global Field CTO, Solo.io

LinkedIn: [/in/cepusta](https://www.linkedin.com/in/cepusta)

X: [@christianposta](https://twitter.com/christianposta)

Github: [christian-posta](https://github.com/christian-posta)

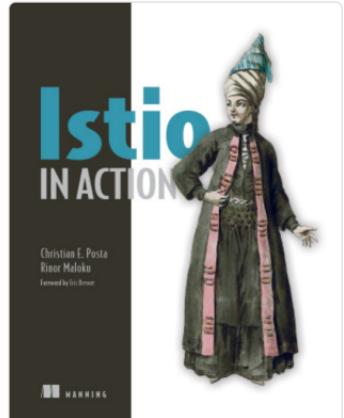
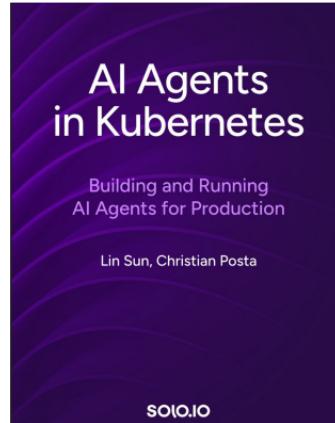
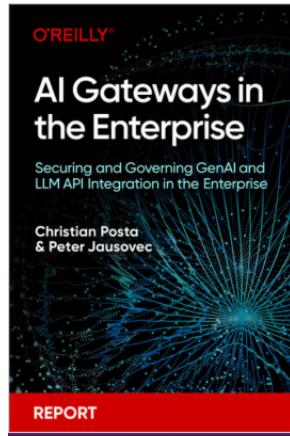


Christian Posta,  
Global Field CTO, Solo.io

LinkedIn: [/in/cepusta](https://www.linkedin.com/in/cepusta)

X: [@christianposta](https://twitter.com/christianposta)

GitHub: [christian-posta](https://github.com/christian-posta)



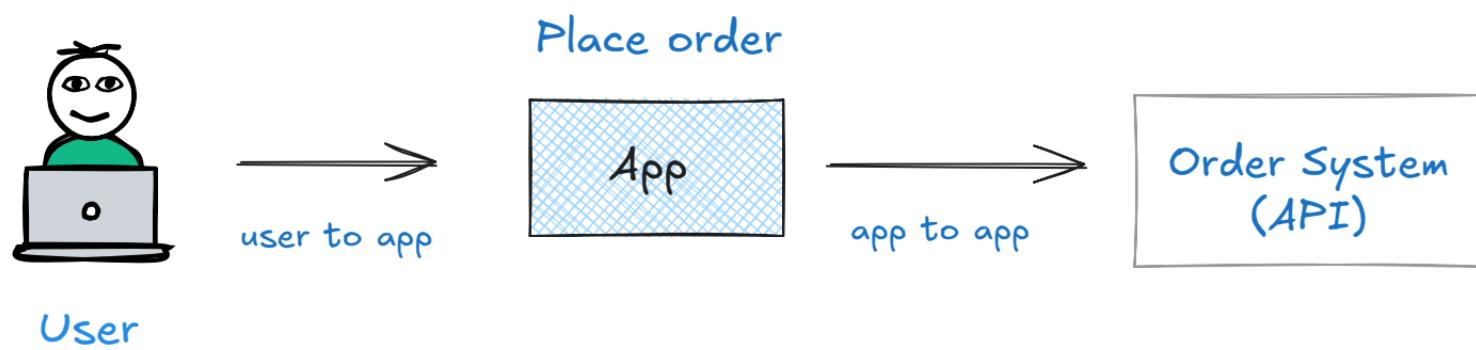
<https://blog.christianposta.com>

**Solo.io**

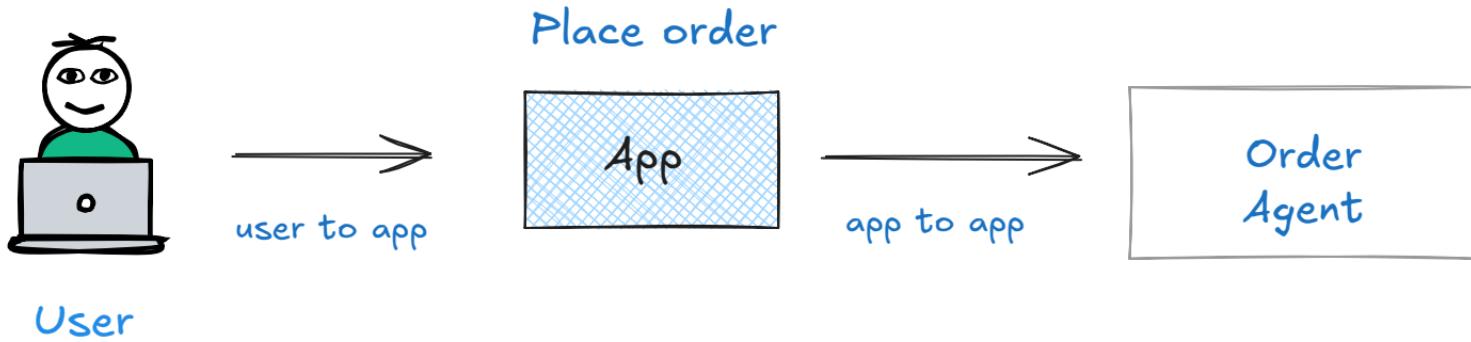
# Agenda

- \* Why Agents and MCP need strong identity
- \* Challenges with existing mechanisms
- \* Technologies that can help
- \* Agent and MCP auth patterns
- \* Demos, demos, and more demos

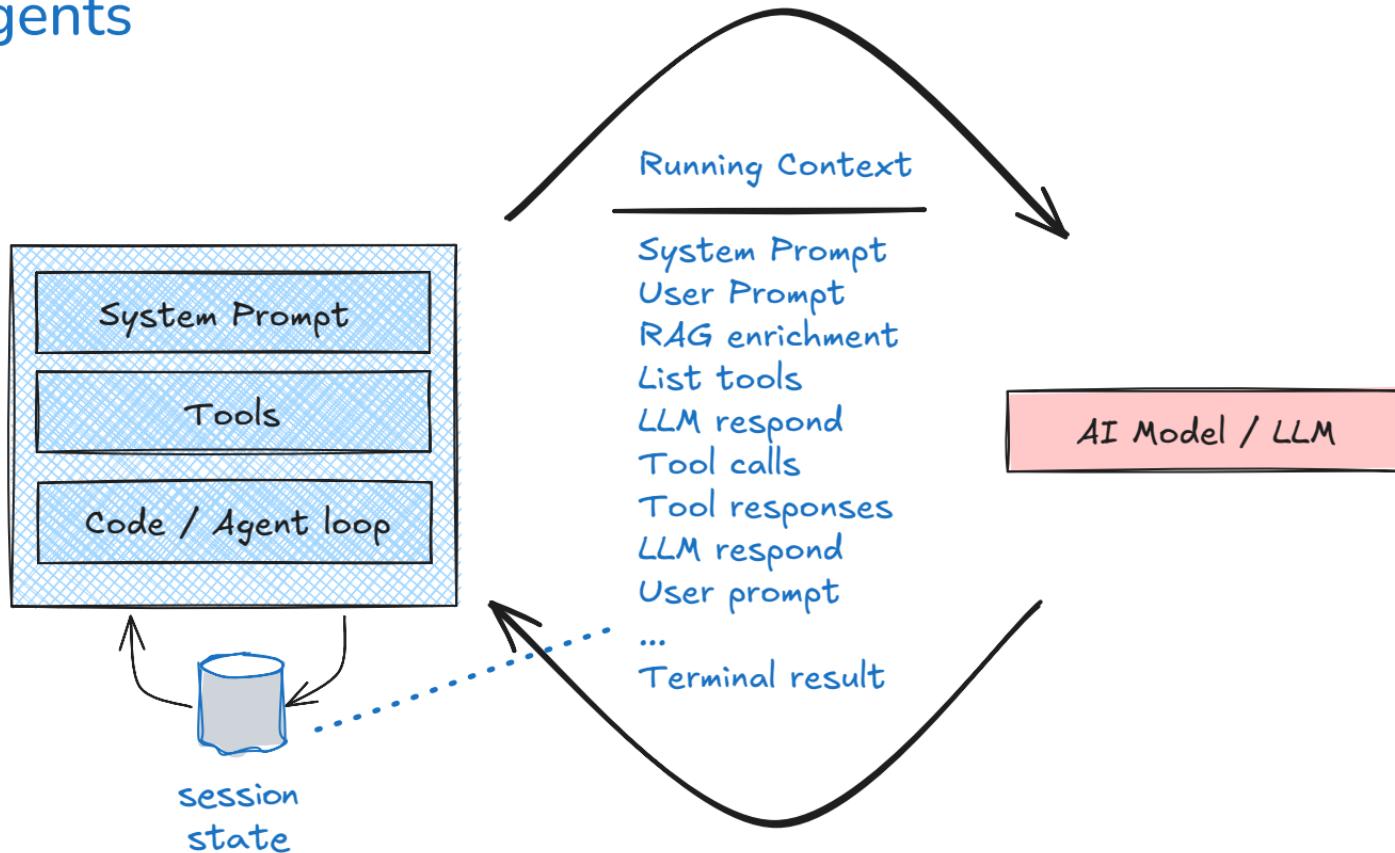
# Typical Application Interaction



# Are Agents Just APIs?



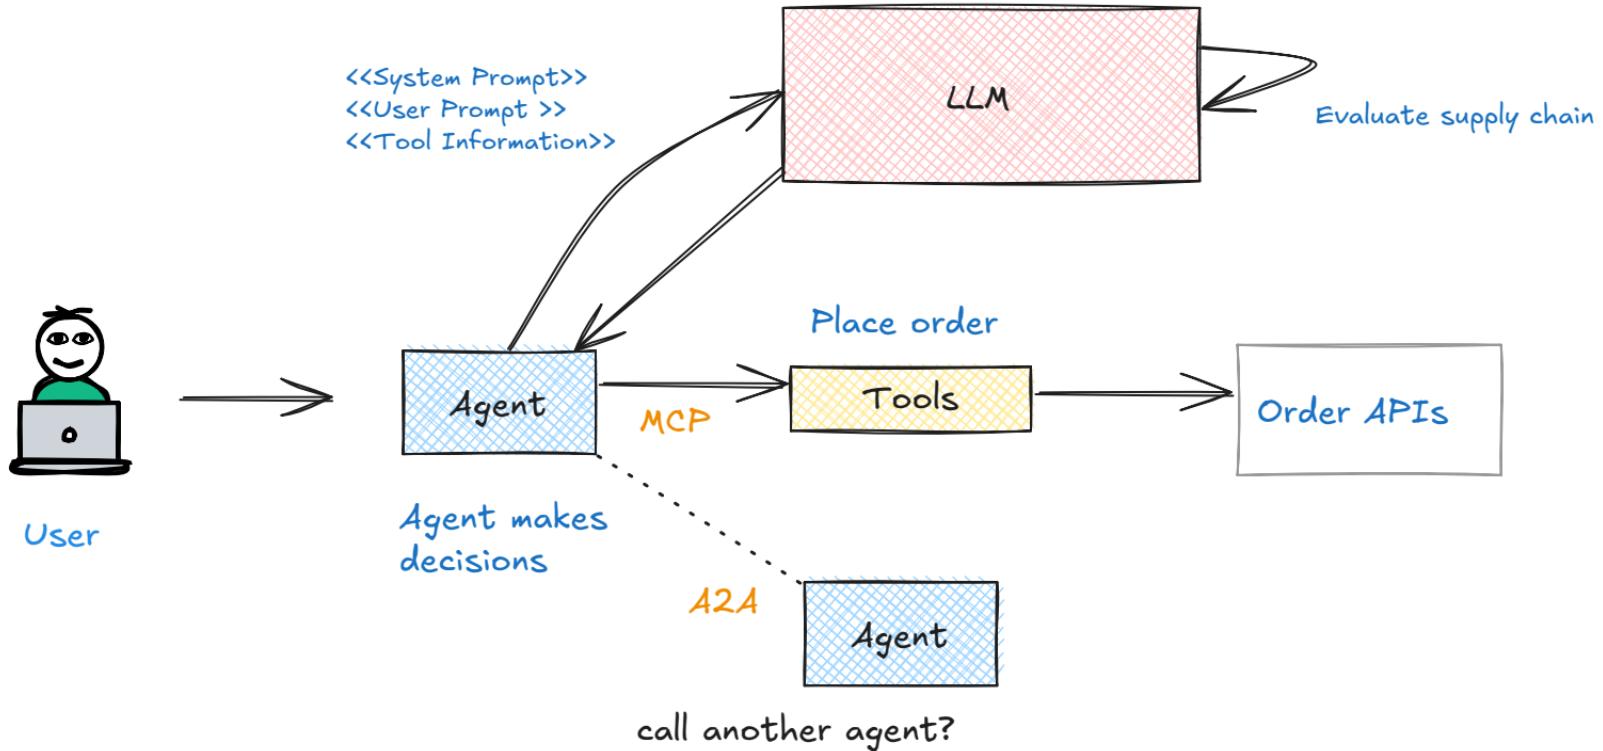
# AI Agents



## Difference Between Agents and APIs

- Agents are stateful, hold sessions
- Agents make decisions based on context
- Agents may act autonomously
- Agents behave non-deterministically
- Agents are goal-driven
- Agents are driven by natural language

# AI Agents



# Auth for MCP and Agents

## Auth Mechanisms Per Spec



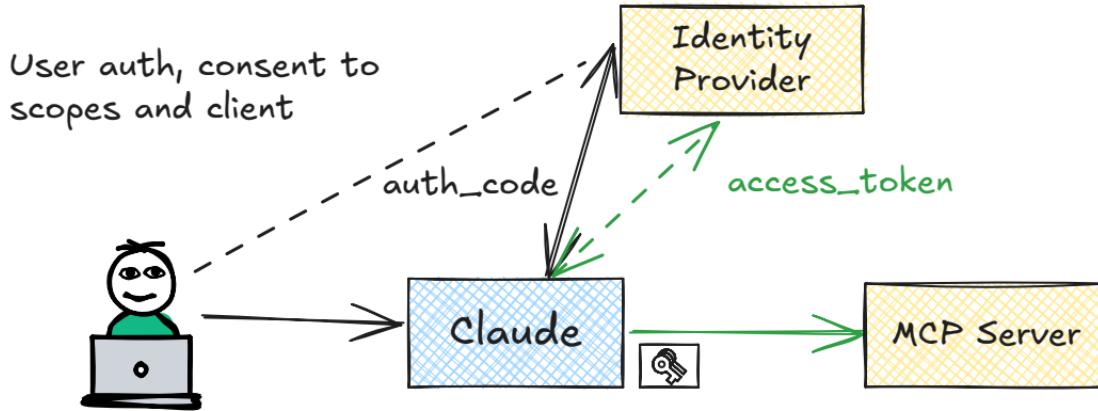
- OAuth 2.1



A2A

- OAuth 2.0 / OIDC
- API Key
- mTLS
- HTTP Basic

# OAuth



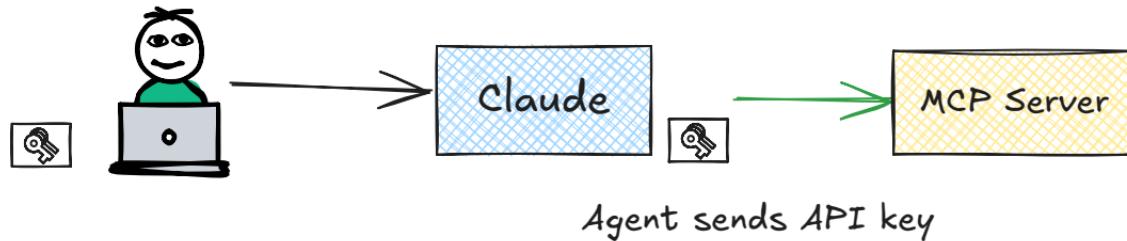
## OAuth for Enterprise?

- enterprise knows users, tools, agents
- "bearer" token semantics
- scopes are too broad
- permissions are static
- no context, intent
- enterprise cares about policy



# API Key

User acquires an API key  
configures API key on Agent



API Key for enterprise?

- API keys were not intended for Auth
- "bearer" token semantics, dangerous
- permissions tied to static RBAC
- usually long-lived, rarely rotated



## Distill Down the Real Enterprise Auth Challenges

- Coarse grained permission models
- Bearer tokens, no proof of possession
- Static / long lived credentials
- No context / intent / capability binding  
(e.g., tenant, data class, risk, etc.).
- Policy bolted on, in code, very inconsistent

# Agent Permissions Example

System / API	Role	Key Permissions
 Oracle HCM	Integration Specialist	<ul style="list-style-type: none"><li>• Read employee personal info</li><li>• View compensation data</li><li>• Retrieve org structures</li><li>• View benefit enrollments</li></ul>
 OpenText Content Server	Integration User	<ul style="list-style-type: none"><li>• See (Level 1)</li><li>• See Contents (Level 2)</li><li>• Read (Level 3)</li><li>• Add (Level 4)</li><li>• Delete (Level 5)</li></ul>
 ServiceNow ITSM	Integration User	<ul style="list-style-type: none"><li>• Create, read, update incidents</li><li>• Assign to self/others</li><li>• Categorize/prioritize</li><li>• Resolve/close incidents</li></ul>

# Agent Permissions Example

👤 **Employee request:** “What’s our company’s work from home policy?”

🤖 **Agent action:** Uses the *Document Management System* API key to retrieve the current WFH policy document from the HR policies folder.

👤 **Employee request:** “I need to update my home address.”

🤖 **Agent action:** Uses the *Ticketing System* API key to create a ticket for the address change and the *Document Management System* API key to retrieve the appropriate form.

👤 **Employee request:** “How many vacation days do I have left this year?”

🤖 **Agent action:** Uses the *HR Management System* API key to query only that specific employee’s time-off balance.

# Agent Permissions Example

## Unauthorized Compensation Access

 **Employee request:** “What’s the average salary for people in my position?”

 **Agent action:** Uses the *HR Management System* API key to query salary data across the entire department or company, accessing compensation information for multiple employees.

 **Why it’s problematic:** The HR Service Account API key has blanket read permissions to all compensation data, allowing the agent to retrieve and analyze sensitive salary information that would normally require manager-level access.

## Sensitive Document Exposure

 **Employee request:** “Are there any reorganization plans for my department?”

 **Agent action:** Uses the *Document Management System* API key to search all HR documents, including confidential executive planning documents labeled as “reorganization.”

 **Why it’s problematic:** The HR Documents Service Role API key can access highly confidential documents that aren’t meant for general employee consumption, potentially revealing sensitive business plans.

## Principals of Agent IAM

- Strong identity for users and agents
- JIT, context/capability tokens issued via policy, bound to identities, agents
- Short lived, frequently rotated
- Enforced by policy evaluation using context
- Least privilege / zero trust posture
- Organized in central systems that can be audited, applied consistently

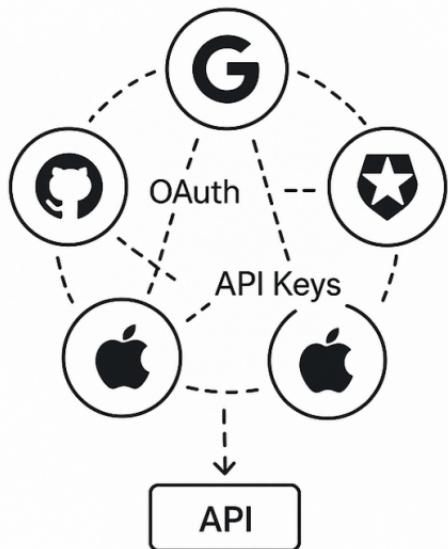
# Laying the Foundation

# Guidance for Enterprise Auth Patterns

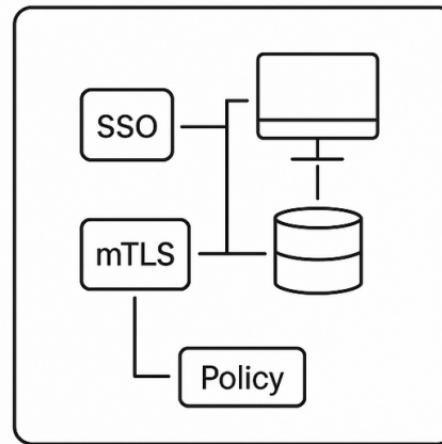
- Tie into enterprise/federated SSO for user auth
- Establish non-human workload/agent identity
- Use mTLS for mutual authentication everywhere
- Use a delegation / on-behalf-of mechanism
- Use centralized policy-driven authorization  
(ABAC/ReBAC)

# Enterprise vs Public

Public Internet



Internal Enterprise



# Identity Based on Attestation



<https://spiffe.io>

SPIFFE ID:

spiffe://acme.com/finance/finance-reporting-agent



trust domain

custom path

SPIFFE Verifiable Document (SVID):



# Secure mTLS Tunnels with SPIFFE

spiffe://acme.com/supply-chain-agent

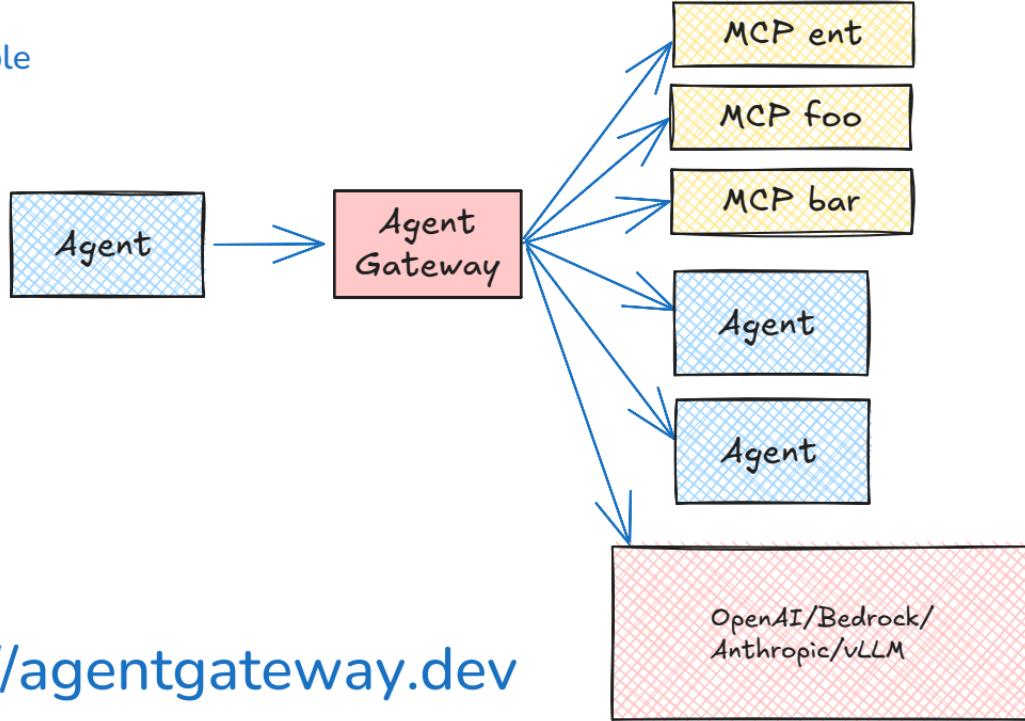


spiffe://acme.com/market-analysis

<https://ambientmesh.io>

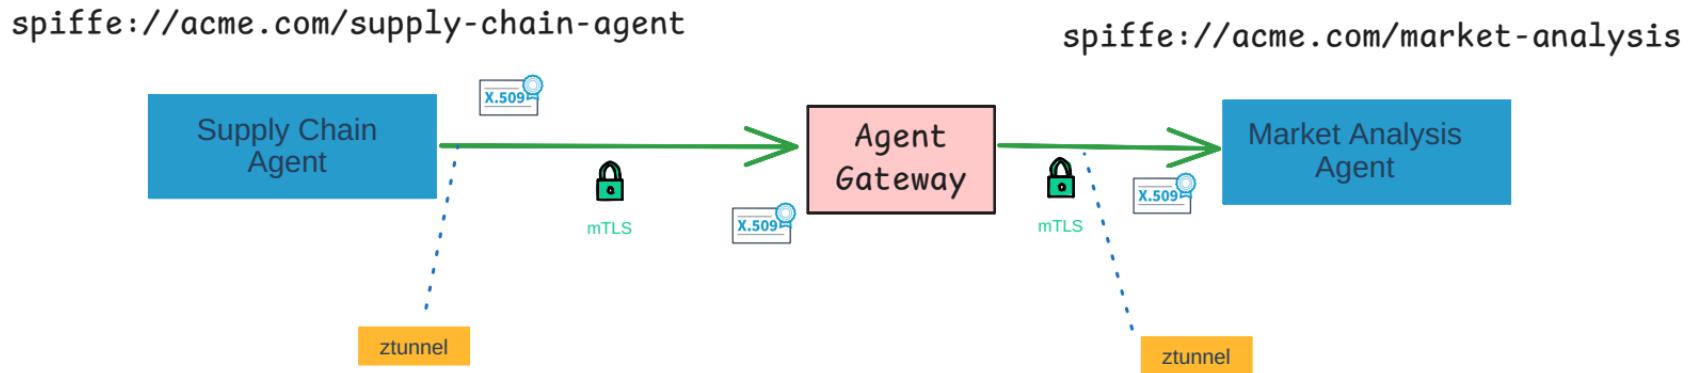
# Agent Gateway for LLM, MCP, and A2A

- \* Kubernetes Gateway API compatible
- \* MCP multiplexing / proxying
- \* A2A routing
- \* JWT handling
- \* Downscope tokens
- \* Auth (user + identity)
- \* Call out to policy engines
- \* Unified API for LLM Calls
- \* Inference gateway routing



<https://agentgateway.dev>

# Agent Gateway with Ambient Mesh

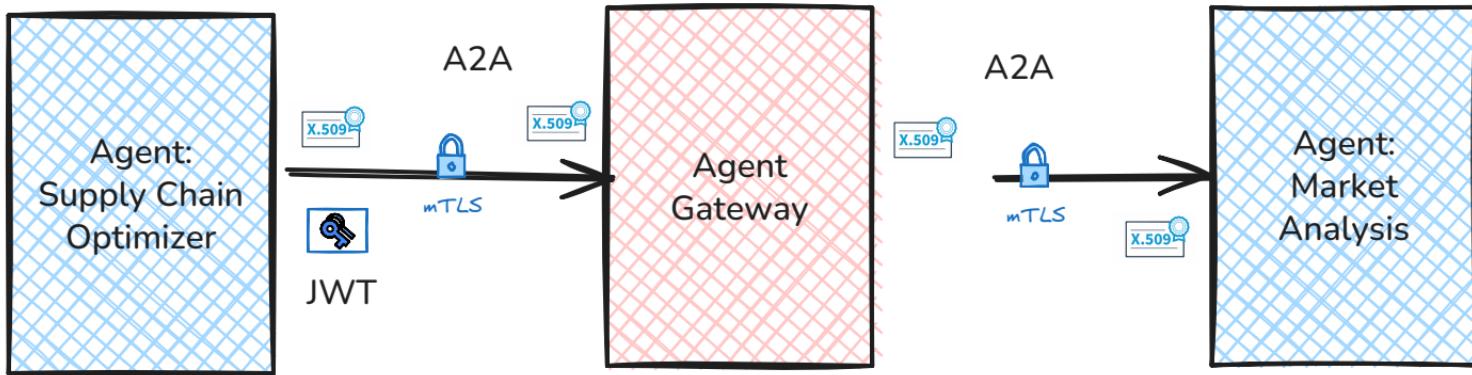


# Auth Patterns

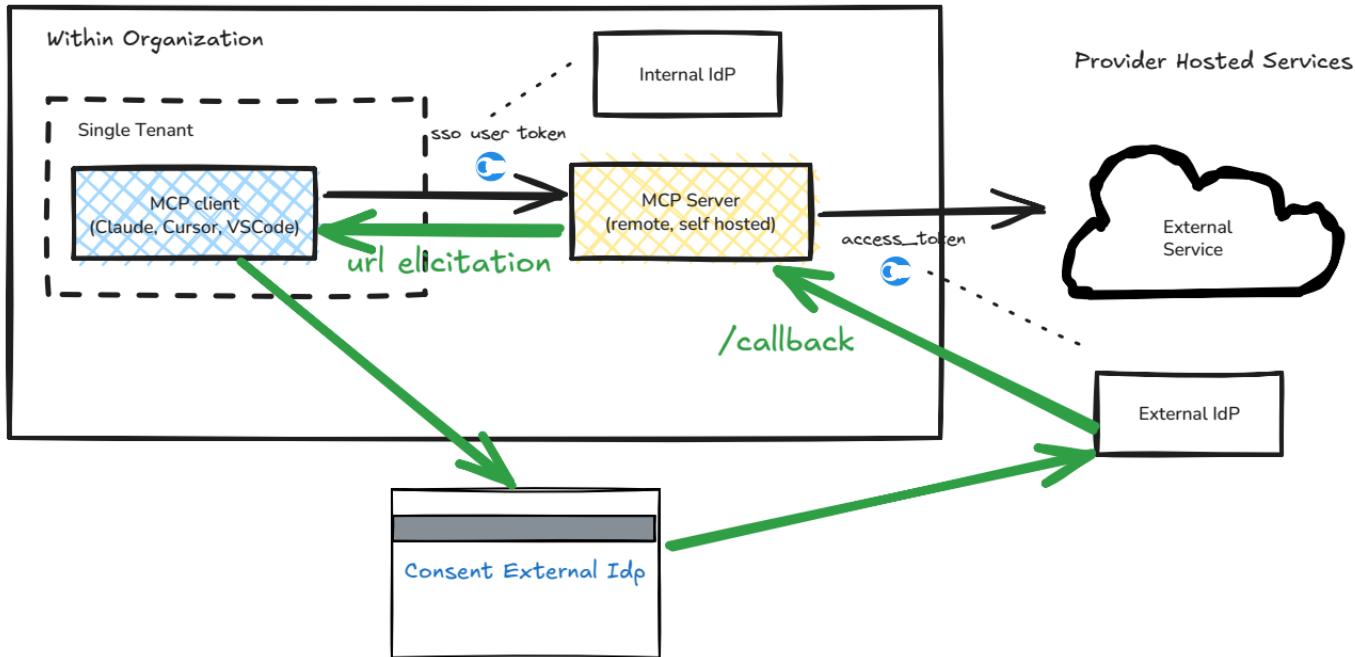
# Patterns for MCP and AI Agents

- Agent to Agent (Internal)
- Agent to MCP server (Internal)
- Agent to MCP (internal) --> API (external)
- Agent (external) to MCP (internal)

# Agent to Agent (internal)

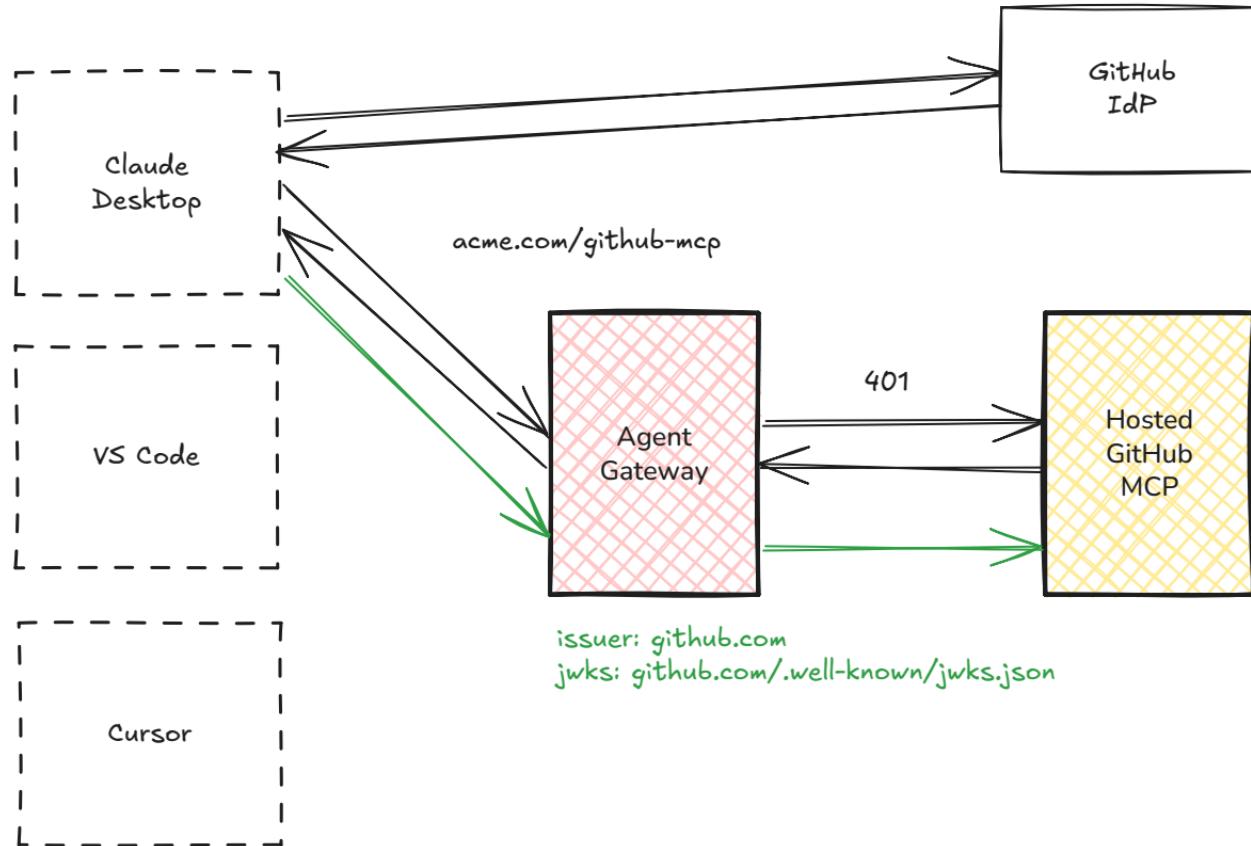






# Agent to MCP (internal) to API (external)

# Agent External to MCP (internal)



Demo

Solo.io

# Thank you!

## Additional Resources:

- \* [ambientmesh.io](https://ambientmesh.io)
- \* [agentgateway.dev](https://agentgateway.dev)
- \* [kagent.dev](https://kagent.dev)
- \* [kgateway.dev](https://kgateway.dev)
- \* [solo.io/blog](https://solo.io/blog)
- \* [blog.christianposta.com](https://blog.christianposta.com)



Christian Posta  
LinkedIn: /in/ceposta  
X: @christianposta  
GitHub: christian-posta

**SOLIO.IO**